Introduction
000

Bayesian Probability
000000000000

Frequentist Statistics
000

Combinations and Permutations
000000000

Distributions
0000000000

Lecture 5 — Probability

Henry Watson

Georgetown University

8/18/23

Morning challenge!

1. Calculate the area under the curve:
   1. $\int_1^5 (8x+3)(4x^2+3x+1)^2$
2. Calculate the gradient and Hessian of the following function:
   1. $f(x,y,z) = 2x + 4xy + 5y^3z^2 + 6x^3y^4$
3. For $x = (1,2,3,4,5)$ and $y = (1.5,4,4,9,14)$
   1. Calculate β and α and for the OLS regression line using the formulas:
      $\hat{\beta} = \frac{\sum_{i=1}^n (X_i - \overline{X})(Y_i - \overline{Y}))}{\sum_{i=1}^n (X_i - \overline{X})^2}$
      $\hat{\alpha} = \overline{Y} - \hat{\beta}\overline{X}$

Agenda

① Introduction

② Bayesian Probability

③ Frequentist Statistics

④ Combinations and Permutations

⑤ Distributions

Frequentist vs. Bayesian

- There are two schools of thought in statistics: Frequentist and Bayesian
- **Frequentist**: The parameters of interest is constant but unknown
    - There is a truth, a right answer, and I can be **confident** to some degree than I have found the truth
- **Bayesian**: The parameters of interest is a random variable.
    - I have an opinion, which I will update based on data to form a new opinion
- Frequentists care about truth, whereas Bayesians care about prediction and perspective
- In academia, we tend to focus more on Frequentist statistics as it is more relevant to causal identification

Frequentist vs. Bayesian

- If I am going to flip a fair coin, what is the probability that it will land heads-side up?
    - Both a frequentist and a bayesian will tell you 50%
- If I have already flipped a coin, but not looked at it yet, what is the probability that that specific coin is heads-side up?
    - The bayesian will still say 50%, but the frequentist will say it's no longer a matter of probability

Definitions:

- **Outcomes**: anything that might happen in the world
- **Events**: composed of one or more outcomes
- **Sample Space**: the set of all possible outcomes
- **Random Events**: events that are probabilistic (as opposed to deterministic
  - When we say "random" here, we mean probabilistic
  - Can identify causal processes that alter the probability, but not causal processes that guarantee the event will occur

Calculating Probability:

$$Pr(e) = \frac{\text{No. of outcomes in event e}}{\text{No. of outcomes in sample space}} \quad (1)$$

e.g. What is the probability of getting a tail when you flip a coin?

Introduction
ooo
Bayesian Probability
o●oooooooooooo
Frequentist Statistics
ooo
Combinations and Permutations
oooooooooo
Distributions
oooooooooo

Relationships between events

- **Independence:** The probability of one event does not change the probability of another event
- **Mutual exclusivity:** If one events occurs, the other event cannot occur
- **Collective Exhaustivity:** Every possible event fits into one of the categories
- **Conditional events::** The probability of one event occurring is affected by whether another event occurs

Bayes theorem:

The *economist* (2004) offers the following illustration of Bayes' rules:
*The canonical example is to imagine that a precarious newborn observes his
first sunrise, and wonders whether the sun will rise again or not. He assigns
equal prior probability to both possible outcomes, and represents this by
placing one white and one black marble into a bag. The following day,
when the sun rises, the child places another white marble in the bag. The
probability that a marble plucked randomly from the bag will be white (i.e.,
the child's degree of belief in future sunrises) has thus gone from a half to
two-thirds. After sunrise the next day, the child adds another white marble,
and the probability (and thus the degree of belief) goes from two-thirds to
three-quarters. And so on. Gradually, the initial belief that the sun is just
as likely to rise as not to rise each morning is modified to become a
near-certainty that sun will always rise.*

Notation!

- $Pr(A)$: the unconditional probability that an event A occurs
  - All probabilities lie between zero and one, so $Pr(A) \in [0; 1]$
- $Pr(A|B)$: the conditional probability of A given B
  - If A is independent of B, $Pr(A|B) = Pr(A)$
- $A \cup B$: the compound event where either A or B happens, or both
- $A \cap B$: the compound event where both A and B happen

More on notation and definitions

- $Pr(A \cap B) = Pr(B|A)Pr(A) = Pr(A|B)Pr(B)$
    - The probability of A and B happening is equal to the conditional probability of B given A times the unconditional probability of A (or vice versa)
    - The probability of someone voting and donating to a candidate is equal to the probability of a donor voting times the probability of donating
- When A and B are independent: $Pr(A \cap B) = Pr(A)Pr(B)$
- Example:
    - Pr(voting) = 0.5
    - Pr(donating) = 0.2
    - If voting is independent of donating, Pr(voting $\cap$ donating) = $0.5 \times 0.2 = 0.10$
    - If Pr(voting|donating) = 0.9, Pr(voting $\cap$ donating) = $0.9 \times 0.2 = 0.18$

More on notation and definition

- When A and B are mutually exclusive: $Pr(A \cap B) = 0$ because $Pr(A|B) = Pr(B|A) = 0$.
- $A \cup B = Pr(A) + Pr(B) - Pr(A \cap B)$
    - The probability of A or B happening is equal to the sum of their unconditional probabilities minus the probability of both happening
    - Using our example from before, Pr(voting $\cup$ donating) $= 0.5 + 0.2 - 0.18 = 0.52$
- We will define $\sim A$ as "not A".

Introduction
000

Bayesian Probability
0000000●000000

Frequentist Statistics
000

Combinations and Permutations
000000000

Distributions
0000000000

Subtract the intersection so you don't double-count it

Bayes theorem:

$$Pr(B|A) = \frac{Pr(A|B)Pr(B)}{P(A)} \tag{2}$$

What is equivalent to:

$$Pr(B|A) = \frac{Pr(A|B)Pr(B)}{P(A|B)Pr(B) + Pr(A|\sim B)Pr(\sim B)} \tag{3}$$

Bayes theorem

Likelihood
How probable is the evidence
given that our hypothesis is true?

Prior
How probable was our hypothesis
before observing the evidence?

$$P(H \mid e) = \frac{P(e \mid H) \, P(H)}{P(e)}$$

Posterior
How probable is our hypothesis
given the observed evidence?
(Not directly computable)

Marginal
How probable is the new evidence
under all possible hypotheses?
$P(e) = \sum P(e \mid H_i) \, P(H_i)$

Introduction
000

Bayesian Probability
0000000000●000

Frequentist Statistics
000

Combinations and Permutations
000000000

Distributions
0000000000

Bayes theorem example

- Covid tests can have false positive (say you're sick when you aren't) and false negative (say you're not sick when you are) results
- 1,000 people have symptoms, and take a test. 50 of them are actually sick.
- 45 people test positive, but 5 of those are false positives
- 955 people test negative, but 10 of those are false negatives

Bayes theorem example

- We can use Bayes Theorem to determine the probability of testing positive if you are sick
- $Pr(Positive|Sick) = \frac{Pr(Sick|Positive)Pr(Positive)}{P(Sick)}$
- $Pr(Sick|Positive) = 40/45 = 0.89$
- $Pr(Positive) = 45/1000 = 0.045$
- $Pr(Sick) = 50/1000 = 0.050$
- $Pr(Positive|Sick) = \frac{(0.89)(0.045)}{0.050} = 0.80$
- **Main lesson:** $Pr(Positive|Sick) \neq Pr(Sick|Positive)$

Odds

The odds of an event happening are defined as the ratio of the probability of the event occurring divided by the probability of the event not occurring. Example: The odds of rolling a four on a dice are $1 : 5$.

$$Odds = \frac{Pr(A)}{1 - Pr(A)} \tag{4}$$

Introduction
○○○
Bayesian Probability
○○○○○○○○○○○○●
Frequentist Statistics
○○○
Combinations and Permutations
○○○○○○○○○
Distributions
○○○○○○○○○○

Odds Ratios

- Odds ratios: the ratio of the odds of an event occurring under the condition A and the odds of that event occurring under the condition B

$$\text{Odds ratio} = \frac{\#Events(A)/\#NonEvents(A)}{\#Events(B)/\#NonEvents(B)} \tag{5}$$

	Small Landlord	Large Landlord
Received Eviction Filing	4,890	13,040
Did not Receive Eviction Filing	86,390	58,860

- Do larger landlords evict at a higher rate?
- $\frac{13040/58860}{4890/86390} = 3.91$

Introduction
000

Bayesian Probability
0000000000000

Frequentist Statistics
●00

Combinations and Permutations
000000000

Distributions
0000000000

Frequentist Methods

- Frequentist statistics relies on parameters of interest being an unknown fixed quantity

- Our methods are built around estimating that truth, and expressing uncertainty

- Examples:
    - A true share of people support Donald Trump for president, and we can estimate that share with a poll
    - That poll will have a "margin of error"
    - There is a true relationship between Democracy and Conflict, and we can estimate that relationship with tools such as linear regression
    - Our estimates will have a "confidence interval"
    - The difference in vote choice between registered Democrats and registered Republicans is *statistically significant*, in that it was unlikely to have occurred by chance

Frequentist Uncertainty

- What is a confidence interval?
- In straightforward terms, a margin of error around an estimate of a parameter which expresses uncertainty
- "Were this procedure to be repeated on numerous samples, the proportion of calculated 95% confidence intervals that encompassed the true value of the population parameter would tend toward 95%."
- We are aiming at an unknown, but constant, parameter. We'll rarely hit it straight on, but if we try many times, our 95% confidence interval will contain the truth 95% of the time. 5% of the time, we'll miss completely, and our confidence interval won't contain the truth.
    - Confidence intervals do not have to be set at 95%, but that is conventional
- A confidence interval should *NOT* be interpreted as "there is a 95% chance that this particular interval contains the true parameter"
    - We live in a frequentist world with constant but unknown parameters; a specific confidence interval either contains the truth or it doesn't

Introduction
ooo

Bayesian Probability
oooooooooooooo

Frequentist Statistics
oo●

Combinations and Permutations
ooooooooo

Distributions
oooooooooo

Confidence Interval Demonstration in R

Combinations and permutations

Combinations: Choose k objects from a set of n objects when the order **does not** matter.

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

Permutations: Choose k objects from a set of n objects when the order **does** matter.

$$\binom{n}{k} = \frac{n!}{(n-k)!}$$

By definition, there will be more permutations that combinations.

Review of Factorials

- $n! = n \times (n-1) \times \cdots \times 1$
- Example: $3! = 3 \times 2 \times 1 = 6$
- Meaning: possible combinations of how a dataset can be arranged
- 1,2,3 ; 1,3,2 ; 2,1,3 ; 2,3,1 ; 3,1,2 ; 3,2,1
  - First number in the combination has three possibilities
  - Second has two possibilities (after the first has been chosen)
  - Third has one possibility (after the first two have been chosen)
- $0! = 1$ (One possible way to arrange a dataset of zero numbers)
- $\frac{n!}{(n-k)!} = n \times (n-1) \cdots \times (n-k+1)$
  - $\frac{5!}{(5-3)!} = 5 \times 4 \times 3$
  - $\frac{5 \times 4 \times 3 \times 2 \times 1}{2 \times 1}$

Why do these formulas work?

Permutations: Choose k objects from a set of n objects when the order **does** matter.

$$\binom{n}{k} = \frac{n!}{(n-k)!}$$

- $n!$ gives you the total number of ways you can arrange all of your objects
- Dividing by $(n-k)!$ reduces the number of permutations based on how many objects you're choosing.
- We only need to consider the first k elements of the factorial in the numerator, because you are only choosing k objects from the set n

Why do these formulas work?

Combinations: Choose k objects from a set of n objects when the order
does not matter.

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

- The additional $k!$ in the denominator eliminates permutations that are
  simply re-orderings of other permutations
- Logic of this is that you can arrange a set of k numbers in $k!$ factorial
  ways
- So if we divide by $k!$, we're left with only one permutation of each
  original combination

Combinations Example:

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

- How many ways can I make a team of 3 out of 5 people?
- $\frac{5!}{3!(5-3)!}$
- $\frac{120}{6(2)}$
- 10

Combinations Example:

- $\frac{5!}{3!(5-3)!}$
- 5! is the number of ways we can order 5 people (120)
- Dividing by $(5-3)!$ reduces this numerator to $5 \times 4 \times 3$, because we only want combinations of three people: 60 permutations
- Additionally dividing by 3! eliminates duplicated (reordered) combinations
    - Example: I can arrange the numbers 3,4,5 in 3! (6) different ways: 3,4,5 ; 3,5,4 ; 4,3,5 ; 4,5,3 ; 5,3,4 ; 5,4,3
    - But these are not original combinations; we have six times as many permutations as unique combinations
    - Dividing by 3! (6) solves this: 10 unique combinations

Combinations Example:

A, B, C, D, E

1. A, B, C
2. A, B, D
3. A, B, E
4. A, C, D
5. A, C, E
6. A, D, E
7. B, C, D
8. B, C, E
9. B, D, E
10. C, D, E

Combinations Example:

12 PhD students decide to host a thanksgiving potluck:

- 3 students will be asked to bring appetizers
- 5 students will be asked to bring main courses
- 3 students will be asked to bring desserts
- The remaining student will be asked to bring drinks

How many different ways can the students be divided into these groups?

$$\binom{12}{3} \cdot \binom{9}{5} \cdot \binom{4}{3} \cdot \binom{1}{1} = \frac{12!}{3!9!} \cdot \frac{9!}{5!4!} \cdot \frac{4!}{3!1!} \cdot \frac{1!}{1!0!} = 220 \cdot 125 \cdot 4 \cdot 1 = 110,880$$

Calculating probabilities

- Common problem: Given a known $Pr(A)$, what is the $Pr(A)$ over many attempts/trials/rolls/etc. if each attempt is independent?
- For example, I play a game of chance with a 10% chance of success. If I play 5 times, what is the probability that I win at least once?
- The probability of winning 5 times in a row is intuitive: $0.1 \times 0.1 \times 0.1 \times 0.1 \times 0.1$ or $(0.1)^5 = 0.00001$
- For the probability of winning *at least once*, we consider the probability of **losing** 5 times in a row, and take the inverse
    - $1 - (1 - 0.1)^5 \approx 0.41$
    - After 10 tries: $1 - (1 - 0.1)^10 \approx 0.65$

Random variables

- Distributions determine the probability that a **random variable** takes on any specific value
- A random variable is **realized** when it takes on any specific value from the set of possible values.
- We say that random variables are **stochastic**: Their probability distribution can be analyzed statistically, but cannot be precisely predicted

Discrete versus continuous

- Remember from the first day:
- Discrete: finite, countable
- Continuous: infinite, uncountable

Sample distribution

- **Sample distribution**: representation of the number of cases that take each value in a sample space
- **Frequency Distribution**: The number of observations/cases for each potential value of a variable
- **Relative Frequency Distribution**: a frequency distribution represented as the share of observations/cases

Contingency tables

- A matrix that shows the joint frequency distribution for two variables
- Also known as "cross-tabulations" or "cross-tabs"

	Dog	Cat	Total
Male	42	10	52
Female	9	39	48
Total	51	49	100

Probability distribution

- The function that describes the likelihood of getting any specific value of a random variable
- Probability of any one value (or range of values) is between 0 and 1
- Sum of probabilities for the full range of values equals 1

Probability Mass Function

For discrete values, we use the Probability Mass Function (effectively a bar

Binomial Distribution
$n = 10$, $p = 0.5$

chart   Binomial distribution: ten trials with $p = 0.5$.

Probability Density Function

- For continuous variables, we use
  the Probability Density Function
- To find the likelihood that an
  observation will be between two
  given values, we look at the area
  under the curve (integration)

(PROBABILITY DENSITY FUNCTION OF PINE TREE HEIGHT)

Cumulative Density Function

The integral, or antiderivative, of the Probability Density Function


CDF for the Normal Distribution

Normal/Gaussian Distribution (the bell curve)

Important features:

- This distribution is symmetric around the center.
- Standard deviations are measures of how spread out the observations are: 68% of values are within one SD, 95% are within two, 99.7% are within three.

Very typical in political science!

Other types of distributions:

- Poisson
- Binomial
- Negative binomial
- t distribution
- F distribution
- Exponential distribution
- Gamma distribution
- And many more...